

✓ Introduction

Problem Statement

To build a CNN based model which can accurately detect melanoma. Melanoma is a type of cancer that can be deadly if not detected early. It accounts for 75% of skin cancer deaths. A solution that can evaluate images and alert dermatologists about the presence of melanoma has the potential to reduce a lot of manual effort needed in diagnosis.

About the Dataset

The dataset consists of 2357 images of malignant and benign oncological diseases, which were formed from the International Skin Imaging Collaboration (ISIC). All images were sorted according to the classification taken with ISIC, and all subsets were divided into the same number of images, with the exception of melanomas and moles, whose images are slightly dominant.

The data set contains the following diseases:

- Actinic keratosis
- Basal cell carcinoma
- Dermatofibroma
- Melanoma
- Nevus
- Pigmented benign keratosis
- Seborrheic keratosis
- Squamous cell carcinoma
- Vascular lesion

✓ Preparation

✓ Setup

```
1 !pip install pycodestyle
2 !pip install --index-url https://test.pypi.org/simple/ nbpep8
3 !pip install Augmentor
```

```

Requirement already satisfied: pycodestyle in /Library/Frameworks/Python.framework
Looking in indexes: https://test.pypi.org/simple/
WARNING: Retrying (Retry(total=4, connect=None, read=None, redirect=None, sta
WARNING: Retrying (Retry(total=3, connect=None, read=None, redirect=None, sta
WARNING: Retrying (Retry(total=2, connect=None, read=None, redirect=None, sta
WARNING: Retrying (Retry(total=1, connect=None, read=None, redirect=None, sta
WARNING: Retrying (Retry(total=0, connect=None, read=None, redirect=None, sta
Could not fetch URL https://test.pypi.org/simple/nbpep8/: There was a problem
ERROR: Could not find a version that satisfies the requirement nbpep8 (from v
ERROR: No matching distribution found for nbpep8
Requirement already satisfied: Augmentor in /Library/Frameworks/Python.framework
Requirement already satisfied: Pillow>=5.2.0 in /Library/Frameworks/Python.framework
Requirement already satisfied: tqdm>=4.9.0 in /Library/Frameworks/Python.framework
Requirement already satisfied: numpy>=1.11.0 in /Library/Frameworks/Python.framework

```

Imports

```

1 import os
2 from pathlib import Path
3 import glob
4 import matplotlib.pyplot as plt
5 from skimage import io
6 import Augmentor
7
8 import tensorflow as tf
9 from tensorflow.keras.models import Sequential
10 from tensorflow.keras.layers import Input, Conv2D, MaxPooling2D, Flatten, \
11     Dense, Dropout
12 from tensorflow.keras.preprocessing.image import ImageDataGenerator, \
13     load_img, img_to_array, DirectoryIterator
14 from tensorflow.keras.regularizers import l2
15
16 from tensorflow.keras.layers import Rescaling, Resizing
17 from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint, \
18     ReduceLROnPlateau
19 from tensorflow.keras.optimizers import Adam
20
21 #from nbpep8.nbpep8 import pep8
22
23 # pep8(_ih)

1 # # Load the Drive helper and mount
2 # from google.colab import drive
3
4 # # This will prompt for authorization.
5 # drive.mount('/content/drive')

```

... Return

▼ setup

▼ Subroutines

```
1 def plot_accuracy(history):
2     """
3     Plots the 'accuracy' and 'val_accuracy'
4     :param history: history from the model fitting
5     :return:
6     """
7     plt.figure(figsize=(12, 4))
8
9     plt.subplot(1, 2, 1)
10    plt.plot(history.history['accuracy'], label='Training Accuracy')
11    plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
12    plt.legend()
13    plt.title('Accuracy')
14
15    plt.subplot(1, 2, 2)
16    plt.plot(history.history['loss'], label='Training Loss')
17    plt.plot(history.history['val_loss'], label='Validation Loss')
18    plt.legend()
19    plt.title('Loss')
20
21    plt.show()
22
23
24 def plot_distribution(directory: str):
25     """
26     Calculates the class distribution in the data directory
27     :param directory: dataset directory with class as subdir
28     :return: dict of class:sample-count
29     """
30    print(f"Looking for files under {os.path.abspath(directory)}")
31    class_size = dict()
32    # Iterate through each subdirectory in the given directory
33    for subdir, _, files in os.walk(directory):
34        if files:
35            condition = subdir.split('/')[-1]
36            class_size.update({condition.upper(): len(files)})
37    return class_size
38
39
40 def distribution(directory: str):
41     """
42     Calculates the class distribution in the data directory
43     :param directory: dataset directory with class as subdir
44     :return: dict of class:sample-count
```

```

45     """"
46     class_size = dict()
47     # Iterate through each subdirectory in the given directory
48     for subdir, _, files in os.walk(directory):
49         if files:
50             condition = str(subdir).split('/')[-1]
51             class_size.update({condition.upper(): len(files)})
52     return class_size
53
54
55 # pep8(_ih)

```

▼ Constants

```

1 DATA_PATH = '/content/drive/MyDrive/MLAI_Notebooks/data/ISIC_skin_cancer/'
2 DATA_PATH = '/Users/pchinnas/Learn/AI/ML/Data/golden/Skin cancer ISIC The Inter
3 BATCH_SIZE = 32
4 RANDOM_SEED = 123
5 MAX_EPOCHS_20 = 20
6 MAX_EPOCHS_30 = 30
7 IMG_SIZE = (180, 180)
8 IMG_CH_SHAPE = (180, 180, 3)
9 INPUT_SHAPE = (BATCH_SIZE, 180, 180, 3)
10 AUG_SAMPLE_PER_CLASS_AUGM = 1000
11 AUG_INC_PER_CLASS_KERAS = 350
12 AUTOTUNE = tf.data.experimental.AUTOTUNE
13
14 # pep8(_ih)

```

▼ Dataset Creation

```

1 os.chdir(DATA_PATH)
2 train_dir = Path("./Train")
3 test_dir = Path('./Test')

1 train_ds = tf.keras.preprocessing.image_dataset_from_directory(
2     train_dir,
3     shuffle=True,
4     image_size=IMG_SIZE,
5     batch_size=BATCH_SIZE
6 )
7
8 len(train_ds)
9 class_names = train_ds.class_names

```

Found 8989 files belonging to 9 classes.

```
1 validation_ds = tf.keras.preprocessing.image_dataset_from_directory(
2     test_dir,
3     shuffle=True,
4     image_size=IMG_SIZE,
5     batch_size=BATCH_SIZE
6 )
```

Found 118 files belonging to 9 classes.

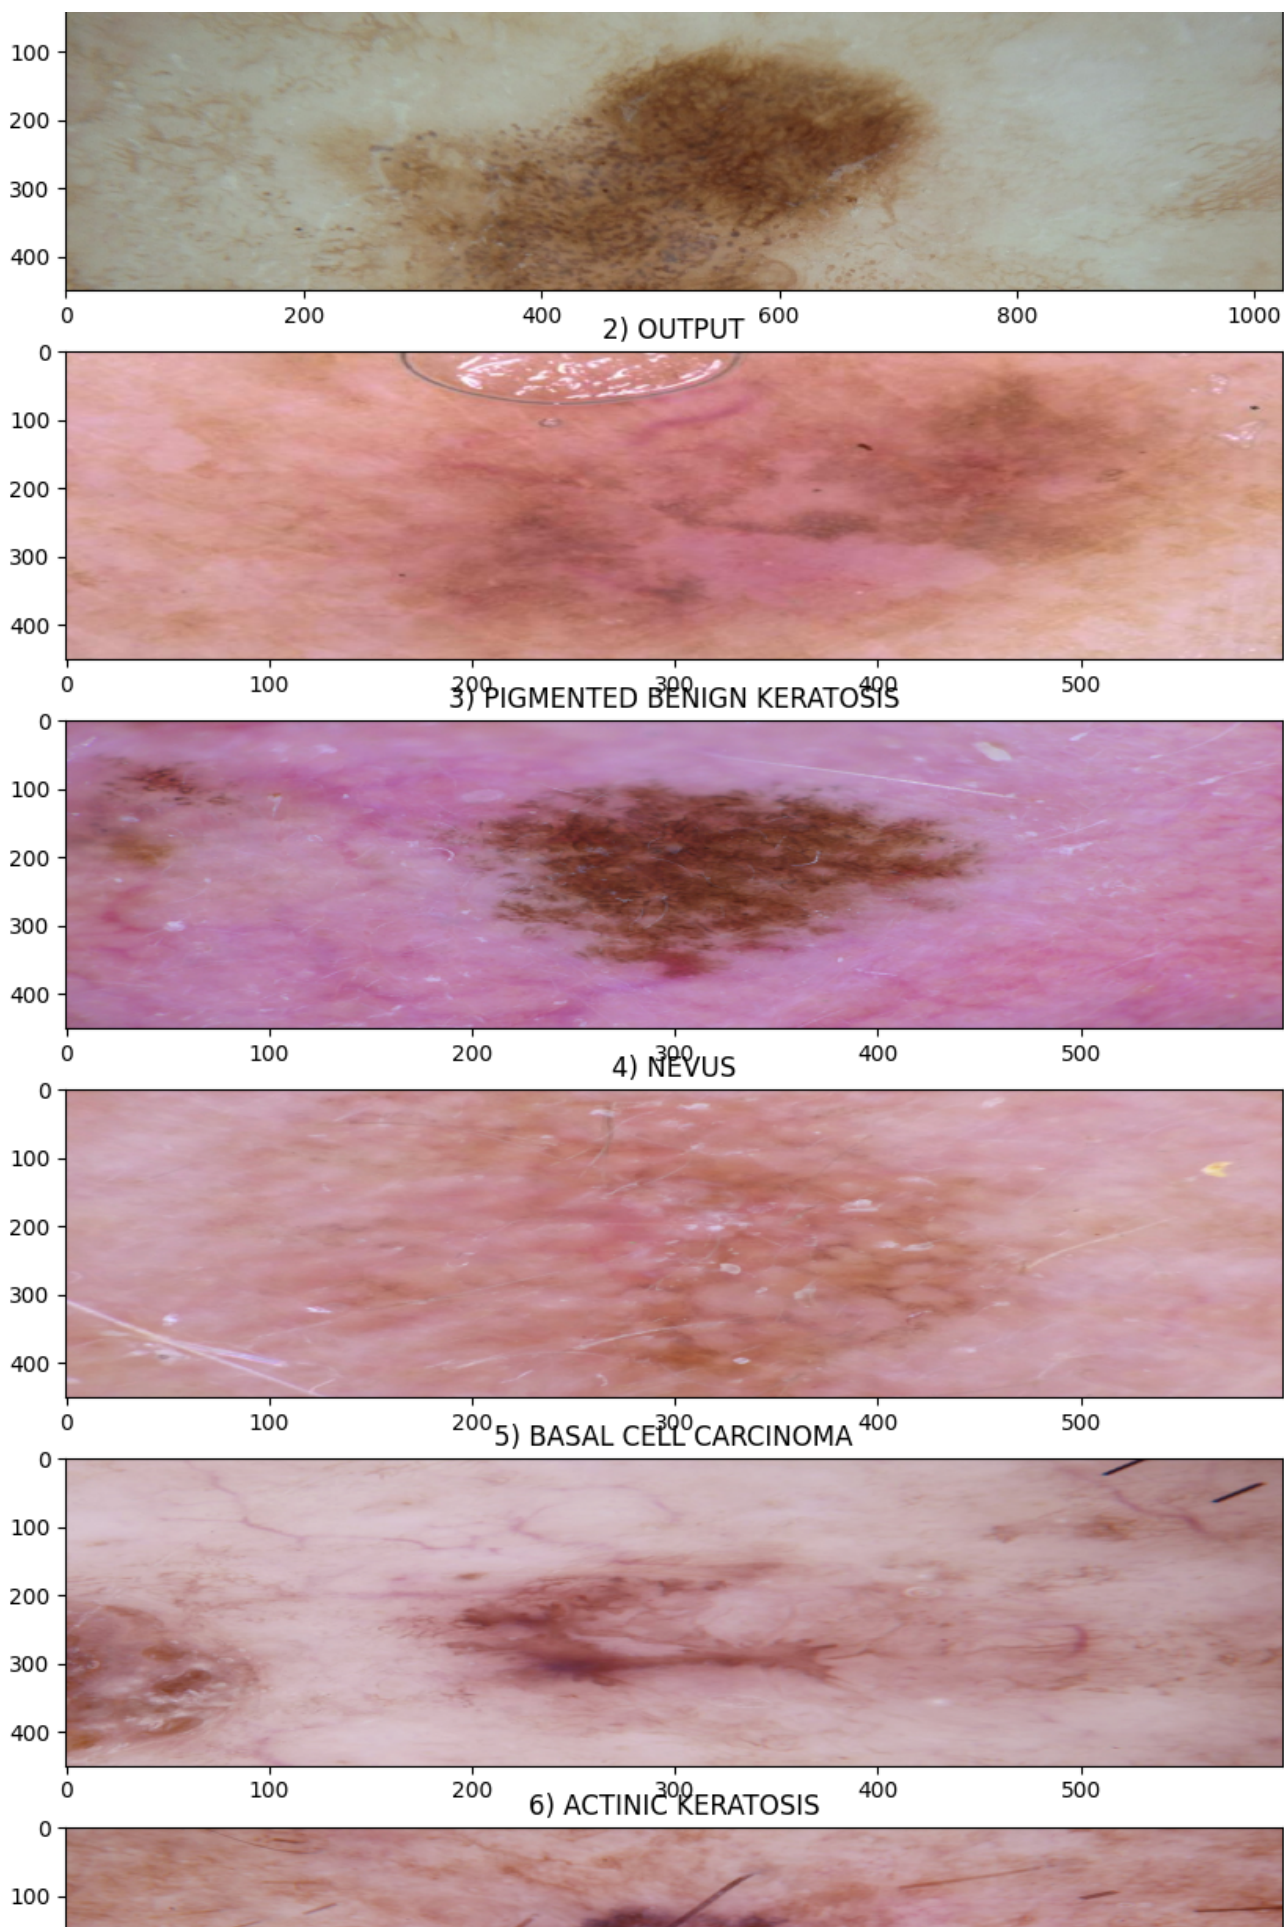
▼ Visualize the Data




```
1 import random
2
3 def pick_cond_img(directory: str):
4     """
5     Picks a random file for each condition
6     :param directory:
7     :return:
8     """
9     selected_files = dict()
10
11     # Iterate through each subdirectory in the given directory
12     for subdir, _, files in os.walk(directory):
13         if files:
14             # Pick a random file from the current subdirectory
15             selected_file = random.choice(files)
16             img_path = os.path.join(subdir, selected_file)
17             condition = str(subdir).split('/')[-1]
18             selected_files.update({condition.upper(): img_path})
19     return selected_files
20 cond_img = pick_cond_img(train_dir)
21
22 f, axes = plt.subplots(len(cond_img), 1, sharey=True, figsize=(10, 30))
23
24 i = 0
25 for cond, img_path in cond_img.items():
26     img = io.imread(img_path)
27     axes[i].imshow(img, aspect='auto')
28     axes[i].title.set_text(str(i+1) + ' ' + cond)
29     i += 1
30 plt.show()
31
32 # pep8(_ih)
```

1) MELANOMA

0







▼ Model - 1


```
1 AUTOTUNE = tf.data.experimental.AUTOTUNE
2 train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
3 validation_ds = validation_ds.cache() \
4     .shuffle(1000).prefetch(buffer_size=AUTOTUNE)
5
6 # pep8(_ih)
```

▼ Class Distribution

```
1 class_dist = distribution(train_dir)
2 plt.pie(class_dist.values(), labels=class_dist.keys(), autopct='%1.1f%%')
3 plt.show()
4 class_dist
5 plot_distribution(os.path.join(DATA_PATH, train_dir))
6
7 # pep8(_ih)
```

OUTPUT





Looking for files under /Users/pchinnas/Learn/AIML/Data/golden/Skin cancer IS:

```
{'MELANOMA': 438,
 'OUTPUT': 750,
 'PIGMENTED BENIGN KERATOSIS': 462,
 'NEVUS': 357,
 'BASAL CELL CARCINOMA': 376,
 'ACTINIC KERATOSIS': 114,
 'SQUAMOUS CELL CARCINOMA': 181,
 'VASCULAR LESION': 139,
 'SEBORRHEIC KERATOSIS': 77,
 'DERMATOFIBROMA': 95}
```

Findings

- 'seborrheic keratosis' has 3.4% (77) of the samples and hence is the least
- 'pigmented benign keratosis' has 20.6%(462) at the highest, closely followed by 'melanoma' at 19.6%(438)

▼ Build

```
1 from tensorflow.keras.layers import Dense, Dropout, Activation, Flatten, \
2     Conv2D, MaxPooling2D, BatchNormalization
3
4 model = Sequential([
5     Input(shape=IMG_CH_SHAPE, batch_size=BATCH_SIZE),
6     Conv2D(32, (3, 3), activation='relu'),
7     Conv2D(64, (3, 3), activation='relu'),
8     MaxPooling2D((2, 2)),
9     Conv2D(64, (3, 3), activation='relu').
```

```

9     conv2d_22 = Conv2D(32, (2, 2), activation='relu',
10     MaxPooling2D((2, 2)),
11     Flatten(),
12     Dense(32, activation='relu'),
13     Dense(64, activation='relu'),
14     Dense(len(class_names), activation='softmax')
15 ])
16
17 # pep8(_ih)

1 model.compile(optimizer='adam',
2               loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
3               metrics=['accuracy'])
4 model.build(INPUT_SHAPE)
5 model.summary()

```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
conv2d_22 (Conv2D)	(32, 178, 178, 32)	896
conv2d_23 (Conv2D)	(32, 176, 176, 64)	18,496
max_pooling2d_20 (MaxPooling2D)	(32, 88, 88, 64)	0
conv2d_24 (Conv2D)	(32, 86, 86, 64)	36,928
max_pooling2d_21 (MaxPooling2D)	(32, 43, 43, 64)	0
flatten_8 (Flatten)	(32, 118336)	0
dense_18 (Dense)	(32, 32)	3,786,784
dense_19 (Dense)	(32, 64)	2,112
dense_20 (Dense)	(32, 9)	585

Total params: 3,845,801 (14.67 MB)
Trainable params: 3,845,801 (14.67 MB)
Non-trainable params: 0 (0.00 B)

▼ Train

```

1 early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
2
3 history = model.fit(
4     train_ds,
5     validation_data=validation_ds,
6     epochs=MAX_EPOCHS_20,
7     callbacks=[early_stopping]
8 )

```

```


9
10 print(f"\nBest accuracy: {max(history.history['accuracy'])}")

Epoch 1/20
/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages
output, from_logits = _get_logits(
281/281 ————— 179s 615ms/step - accuracy: 0.1477 - loss: 14.410
Epoch 2/20
281/281 ————— 181s 642ms/step - accuracy: 0.1929 - loss: 2.131
Epoch 3/20
281/281 ————— 183s 651ms/step - accuracy: 0.2575 - loss: 1.951
Epoch 4/20
281/281 ————— 181s 644ms/step - accuracy: 0.3209 - loss: 1.834
Epoch 5/20
281/281 ————— 187s 667ms/step - accuracy: 0.2050 - loss: 2.060
Epoch 6/20
281/281 ————— 182s 649ms/step - accuracy: 0.4098 - loss: 1.656
Epoch 7/20
281/281 ————— 186s 661ms/step - accuracy: 0.4270 - loss: 1.579

Best accuracy: 0.4707976281642914

```

```
1 plot_accuracy(history)
```



Findings

- Accuracy is low at **47.5%**. This can be attributed to the following
 - Not having enough labeled training data can lead to poor generalization. CNNs

require large amounts of diverse data to learn effectively.

- **Imbalanced Dataset:** If some classes have significantly more samples than others, the model might become biased toward the more frequent classes, leading to poor performance on the less frequent ones.
- **Overfitting:** When a model performs well on training data but poorly on validation data, it might be overfitting. This can occur if the model is too complex relative to the amount of training data.

✓ Model - 2

✓ Build

```
1 # resize and rescale the value
2 resize_and_rescale = Sequential([
3     tf.keras.layers.Resizing(180, 180),
4     tf.keras.layers.Rescaling(1.0/255)
5 ])
```

```
1 keras_augmentation = Sequential([
2     tf.keras.layers.RandomRotation(factor=(-0.2, 0.3)),
3     tf.keras.layers.RandomFlip("horizontal_and_vertical")
4 ])
```

```
1 from tensorflow.keras.layers import Rescaling, Resizing
2
3 model = Sequential([
4     Input(shape=IMG_CH_SHAPE, batch_size=BATCH_SIZE),
5     resize_and_rescale,
6     keras_augmentation,
7     Conv2D(32, (3, 3), activation='relu'),
8     Conv2D(64, (3, 3), activation='relu'),
9     MaxPooling2D((2, 2)),
10    Conv2D(64, (3, 3), activation='relu'),
11    MaxPooling2D((2, 2)),
12    Flatten(),
13    Dense(32, activation='relu'),
14    Dense(64, activation='relu'),
15    Dense(len(class_names), activation='softmax')
16 ])
```

```
1 model.compile(optimizer='adam',
```

```

2         loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True)
3         metrics=['accuracy'])
4 model.build(INPUT_SHAPE)
5 model.summary()

```

Model: "sequential_13"

Layer (type)	Output Shape	Param #
sequential_11 (Sequential)	(32, 180, 180, 3)	0
sequential_12 (Sequential)	(32, 180, 180, 3)	0
conv2d_25 (Conv2D)	(32, 178, 178, 32)	896
conv2d_26 (Conv2D)	(32, 176, 176, 64)	18,496
max_pooling2d_22 (MaxPooling2D)	(32, 88, 88, 64)	0
conv2d_27 (Conv2D)	(32, 86, 86, 64)	36,928
max_pooling2d_23 (MaxPooling2D)	(32, 43, 43, 64)	0
flatten_9 (Flatten)	(32, 118336)	0
dense_21 (Dense)	(32, 32)	3,786,784
dense_22 (Dense)	(32, 64)	2,112
dense_23 (Dense)	(32, 9)	585

Total params: 3,845,801 (14.67 MB)
Trainable params: 3,845,801 (14.67 MB)
Non-trainable params: 0 (0.00 B)


▼ Train

```


1 early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
2
3 history = model.fit(
4     train_ds,
5     validation_data=validation_ds,
6     epochs=MAX_EPOCHS_20,
7     callbacks=[early_stopping]
8 )
9
10 print(f"\nBest accuracy: {max(history.history['accuracy'])}")

```

Epoch 1/20

281/281  **191s** 676ms/step - accuracy: 0.1951 - loss: 2.1151

Epoch 2/20


281/281  **180s** 642ms/step - accuracy: 0.4016 - loss: 1.5511

Epoch 3/20

281/281 ————— 175s 624ms/step - accuracy: 0.4312 - loss: 1.4741
 Epoch 4/20
 281/281 ————— 192s 683ms/step - accuracy: 0.4440 - loss: 1.4180
 Epoch 5/20
 281/281 ————— 192s 684ms/step - accuracy: 0.4685 - loss: 1.3561
 Epoch 6/20
 281/281 ————— 188s 668ms/step - accuracy: 0.4739 - loss: 1.3321

Best accuracy: 0.47547000646591187

1 plot_accuracy(history)



Findings

- Validation accuracy has not improved (**47.5%**)
- There is a clear case of overfitting
- Total trainable params are 3,845,801
- The class imbalance can be a cause of this overfitting
- A good sample size can help the model quality

✓ Model - 3

✓ Augmenting Data


```

1 import Augmentor
2 for class_name in class_names:
3     p = Augmentor.Pipeline(os.path.join(train_dir, class_name))
4     p.rotate(probability=0.7, max_left_rotation=10, max_right_rotation=10)
5     p.sample(AUG_SAMPLE_PER_CLASS_AUGM)
6
7 # pep8(_ih)

    Initialised with 114 image(s) found.
    Output directory set to Train/actinic keratosis/output.Processing <PIL.Image.
    Initialised with 376 image(s) found.
    Output directory set to Train/basal cell carcinoma/output.Processing <PIL.Ima
    Initialised with 95 image(s) found.
    Output directory set to Train/dermatofibroma/output.Processing <PIL.Image.Ima
    Initialised with 438 image(s) found.
    Output directory set to Train/melanoma/output.Processing <PIL.Image.Image ima
    Initialised with 357 image(s) found.
    Output directory set to Train/nevus/output.Processing <PIL.Image.Image image
    Initialised with 462 image(s) found.
    Output directory set to Train/pigmented benign keratosis/output.Processing <P
    Initialised with 77 image(s) found.
    Output directory set to Train/seborrheic keratosis/output.Processing <PIL.Ima
    Initialised with 181 image(s) found.
    Output directory set to Train/squamous cell carcinoma/output.Processing <PIL.
    Initialised with 139 image(s) found.
    Output directory set to Train/vascular lesion/output.Processing <PIL.Image.Ima

```

```

1 train_ds = tf.keras.preprocessing.image_dataset_from_directory(
2     train_dir,
3     seed=RANDOM_SEED,
4     validation_split=0.2,
5     subset='training',
6     image_size=IMG_SIZE,
7     batch_size=BATCH_SIZE
8 )
9
10 # pep8(_ih)

```

```

    Found 11239 files belonging to 9 classes.
    Using 8992 files for training.

```

```

1 validation_ds = tf.keras.preprocessing.image_dataset_from_directory(
2     train_dir,
3     seed=RANDOM_SEED,
4     validation_split=0.2,
5     subset='validation',
6     image_size=IMG_SIZE,
7     batch_size=BATCH_SIZE
8 )
~

```

```
9
10 # pep8(_ih)

Found 11239 files belonging to 9 classes.
Using 2247 files for validation.

1 train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
2 validation_ds = validation_ds.cache().prefetch(buffer_size=AUTOTUNE)
```

▼ Build

```
1 model = Sequential([
2     Input(shape=IMG_CH_SHAPE, batch_size=BATCH_SIZE),
3
4     Rescaling(1.0/255),
5
6     Conv2D(16, (3, 3), activation='relu'),
7     BatchNormalization(),
8     MaxPooling2D((2, 2)),
9
10    Conv2D(32, (3, 3), activation='relu'),
11    BatchNormalization(),
12    MaxPooling2D((2, 2)),
13
14    Conv2D(64, (3, 3), activation='relu', kernel_regularizer=l2(0.01)),
15    BatchNormalization(),
16    MaxPooling2D((2, 2)),
17
18    Conv2D(64, (3, 3), activation='relu', kernel_regularizer=l2(0.01)),
19    BatchNormalization(),
20    MaxPooling2D((2, 2)),
21
22    # Conv2D(64, (3, 3), activation='relu', kernel_regularizer=l2(0.01)),
23    # BatchNormalization(),
24    # MaxPooling2D((2, 2)),
25
26    Dropout(0.25),
27
28    Flatten(),
29
30    Dense(128, activation='relu', kernel_regularizer=l2(0.01)),
31
32    Dropout(0.5),
33
34    Dense(len(class_names), activation='softmax')
35 ])
36
37 # pep8(_ih)
```

```

1 adam = Adam(learning_rate=0.0001)
2
3 model.compile(optimizer=adam,
4               loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=T
5               metrics=['accuracy']))
6 model.build(INPUT_SHAPE)
7 model.summary()
8
9 # pep8(_ih)

```

Model: "sequential_16"

Layer (type)	Output Shape	Param #
rescaling_10 (Rescaling)	(32, 180, 180, 3)	0
conv2d_36 (Conv2D)	(32, 178, 178, 16)	448
batch_normalization_19 (BatchNormalization)	(32, 178, 178, 16)	64
max_pooling2d_32 (MaxPooling2D)	(32, 89, 89, 16)	0
conv2d_37 (Conv2D)	(32, 87, 87, 32)	4,640
batch_normalization_20 (BatchNormalization)	(32, 87, 87, 32)	128
max_pooling2d_33 (MaxPooling2D)	(32, 43, 43, 32)	0
conv2d_38 (Conv2D)	(32, 41, 41, 64)	18,496
batch_normalization_21 (BatchNormalization)	(32, 41, 41, 64)	256
max_pooling2d_34 (MaxPooling2D)	(32, 20, 20, 64)	0
conv2d_39 (Conv2D)	(32, 18, 18, 64)	36,928
batch_normalization_22 (BatchNormalization)	(32, 18, 18, 64)	256
max_pooling2d_35 (MaxPooling2D)	(32, 9, 9, 64)	0
dropout_16 (Dropout)	(32, 9, 9, 64)	0
flatten_12 (Flatten)	(32, 5184)	0
dense_28 (Dense)	(32, 128)	663,680
dropout_17 (Dropout)	(32, 128)	0
dense_29 (Dense)	(32, 9)	1,161

Total params: 726,057 (2.77 MB)

```

total params: 725,705 (2.77 MB)
Trainable params: 725,705 (2.77 MB)
Non-trainable params: 352 (1.38 KB)

```

▼ Train

```

1 optimizer = Adam(learning_rate=0.0001)
2 early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
3 checkpoint = ModelCheckpoint("model.keras", monitor="val_loss", save_best_only=True)
4 reduce_lr_loss = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=7, verbose=1)
5
6 history = model.fit(
7     train_ds,
8     validation_data=validation_ds,
9     epochs=MAX_EPOCHS_30,
10    callbacks=[early_stopping, checkpoint, reduce_lr_loss]
11 )
12 print(f"\nBest accuracy: {max(history.history['accuracy'])}")
13
14 # pep8(_ih)

```



Epoch 1/30
 /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages/keras/callbacks/early_stopping.py:100: UserWarning: EarlyStopping: val_loss did not improve over 5 epochs. Stopping training.
 output, from_logits = _get_logits(
 281/281 ————— 0s 179ms/step - accuracy: 0.2228 - loss: 6.1990
 Epoch 1: val_loss improved from inf to 6.82118, saving model to model.keras
 281/281 ————— 60s 194ms/step - accuracy: 0.2231 - loss: 6.1969
 Epoch 2/30
 281/281 ————— 0s 174ms/step - accuracy: 0.4506 - loss: 4.8954
 Epoch 2: val_loss improved from 6.82118 to 4.71839, saving model to model.keras
 281/281 ————— 51s 183ms/step - accuracy: 0.4507 - loss: 4.8951
 Epoch 3/30
 281/281 ————— 0s 174ms/step - accuracy: 0.5290 - loss: 4.4855
 Epoch 3: val_loss improved from 4.71839 to 3.99648, saving model to model.keras
 281/281 ————— 51s 183ms/step - accuracy: 0.5291 - loss: 4.4853
 Epoch 4/30
 281/281 ————— 0s 173ms/step - accuracy: 0.5857 - loss: 4.1182
 Epoch 4: val_loss improved from 3.99648 to 3.69043, saving model to model.keras
 281/281 ————— 51s 182ms/step - accuracy: 0.5857 - loss: 4.1179
 Epoch 5/30
 281/281 ————— 0s 173ms/step - accuracy: 0.6489 - loss: 3.7791
 Epoch 5: val_loss improved from 3.69043 to 3.42252, saving model to model.keras
 281/281 ————— 51s 182ms/step - accuracy: 0.6489 - loss: 3.7788
 Epoch 6/30
 281/281 ————— 0s 175ms/step - accuracy: 0.6840 - loss: 3.4914
 Epoch 6: val_loss improved from 3.42252 to 3.15235, saving model to model.keras
 281/281 ————— 52s 184ms/step - accuracy: 0.6840 - loss: 3.4912
 Epoch 7/30
 281/281 ————— 0s 176ms/step - accuracy: 0.7282 - loss: 3.1983
 Epoch 7: val_loss improved from 3.15235 to 2.95956, saving model to model.keras
 281/281 ————— 52s 186ms/step - accuracy: 0.7282 - loss: 3.1982
 Epoch 8/30
 281/281 ————— 0s 170ms/step - accuracy: 0.7519 - loss: 2.9637

```

Epoch 8: val_loss improved from 2.95956 to 2.76957, saving model to model.keras
281/281 ━━━━━━━━━━━━━━━━━ 50s 178ms/step - accuracy: 0.7519 - loss: 2.9636
Epoch 9/30
281/281 ━━━━━━━━━━━━━━━━━ 0s 168ms/step - accuracy: 0.7805 - loss: 2.7642
Epoch 9: val_loss improved from 2.76957 to 2.51044, saving model to model.keras
281/281 ━━━━━━━━━━━━━━━━━ 50s 177ms/step - accuracy: 0.7805 - loss: 2.7641
Epoch 10/30
281/281 ━━━━━━━━━━━━━━━━━ 0s 173ms/step - accuracy: 0.8092 - loss: 2.5453
Epoch 10: val_loss improved from 2.51044 to 2.45920, saving model to model.keras
281/281 ━━━━━━━━━━━━━━━━━ 51s 182ms/step - accuracy: 0.8092 - loss: 2.5452
Epoch 11/30
281/281 ━━━━━━━━━━━━━━━━━ 0s 170ms/step - accuracy: 0.8398 - loss: 2.3497
Epoch 11: val_loss improved from 2.45920 to 2.19092, saving model to model.keras
281/281 ━━━━━━━━━━━━━━━━━ 50s 178ms/step - accuracy: 0.8398 - loss: 2.3497
Epoch 12/30
281/281 ━━━━━━━━━━━━━━━━━ 0s 171ms/step - accuracy: 0.8482 - loss: 2.1986
Epoch 12: val_loss improved from 2.19092 to 2.06845, saving model to model.keras
281/281 ━━━━━━━━━━━━━━━━━ 51s 183ms/step - accuracy: 0.8482 - loss: 2.1985
Epoch 13/30
281/281 ━━━━━━━━━━━━━━━━━ 0s 172ms/step - accuracy: 0.8607 - loss: 2.0522
Epoch 13: val_loss improved from 2.06845 to 1.95785, saving model to model.keras
281/281 ━━━━━━━━━━━━━━━━━ 51s 182ms/step - accuracy: 0.8607 - loss: 2.0521
Epoch 14/30
281/281 ━━━━━━━━━━━━━━━━━ 0s 172ms/step - accuracy: 0.8832 - loss: 1.9156
Epoch 14: val_loss improved from 1.95785 to 1.91078, saving model to model.keras
281/281 ━━━━━━━━━━━━━━━━━ 52s 183ms/step - accuracy: 0.8831 - loss: 1.9156

```

```
1 plot_accuracy(history)
```



Findings

- After increasing the size of the training dataset, the accuracy has improved (**94.6%**).
- Total trainable params are 725,705
- Overfitting is contained with the help of
 - BatchNormalization
 - Dropouts
 - EarlyStopping
- The model also runs with a Learning rate scheduler both with ReduceLROnPlateau and with Adam optimizer with less lr value

✓ Analysis

```
1 model.evaluate(validation_ds)
```

```
71/71 ————— 3s 35ms/step - accuracy: 0.8926 - loss: 0.9721
[0.9695644378662109, 0.8985313773155212]
```

```
1 print("Best Score:", max(history.history['accuracy']))
```

```
Best Score: 0.9466192126274109
```

```
1 acc = ['accuracy', 'val_accuracy']
2 loss = ['loss', 'val_loss']
3 for s in acc:
4     print(s, max(history.history[s]))
5 for s in loss:
6     print(s, min(history.history[s]))
7
8 # pep8(_ih)
```

```
accuracy 0.9466192126274109
val_accuracy 0.9092122912406921
loss 0.8464059233665466
val_loss 0.9695644378662109
```

✓ Prediction:

```
1 from glob import glob
2 import numpy as np
3 class_name = class_names[1]
4 Test_image_path = os.path.join(test_dir, class_name, '*')
5 print(Test_image_path)
```



```

5 print(test_image_path)
6 Test_image = glob(Test_image_path)
7 Test_image = load_img(Test_image[-1], target_size=(180, 180, 3))
8 plt.imshow(Test_image)
9 plt.grid(False)
10
11 img = np.expand_dims(Test_image, axis=0)
12 pred = model.predict(img)
13 pred = np.argmax(pred)
14 pred_class = class_names[pred]
15 print(f"Actual Class '{class_name} \nPredictive Class '{pred_class}'")
16
17 # pep8(_ih)


```

Test/basal cell carcinoma/*

1/1 ————— 0s 67ms/step

Actual Class 'basal cell carcinoma'

Predictive Class 'basal cell carcinoma'



✓ Model on Unseen Test data

```

1 test_ds = tf.keras.preprocessing.image_dataset_from_directory(
2     test_dir,
3     seed=RANDOM_SEED,
4     validation_split=0.2,
5     subset='validation',
6     image_size=IMG_SIZE,

```

```

7     batch_size=BATCH_SIZE
8 )
9

```

```

Found 118 files belonging to 9 classes.
Using 23 files for validation.

```

```

1 optimizer = Adam(learning_rate=0.0001)
2 early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_we
3 checkpoint = ModelCheckpoint("model.keras",monitor="val_loss",save_best_only=T
4 reduce_lr_loss = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=7,
5
6 history = model.fit(
7     train_ds,
8     validation_data=test_ds,
9     epochs=MAX_EPOCHS_30,
10    callbacks=[early_stopping, checkpoint, reduce_lr_loss]
11 )
12 print(f"\nBest accuracy: {max(history.history['accuracy'])}")
13
14 # pep8(_ih)

```

```

Epoch 1/30
281/281 ————— 0s 173ms/step - accuracy: 0.9450 - loss: 0.8325
Epoch 1: val_loss improved from inf to 3.87447, saving model to model.keras
281/281 ————— 49s 174ms/step - accuracy: 0.9450 - loss: 0.8325
Epoch 2/30
281/281 ————— 0s 208ms/step - accuracy: 0.9476 - loss: 0.8045
Epoch 2: val_loss did not improve from 3.87447
281/281 ————— 59s 209ms/step - accuracy: 0.9476 - loss: 0.8045
Epoch 3/30
281/281 ————— 0s 164ms/step - accuracy: 0.9447 - loss: 0.7874
Epoch 3: val_loss did not improve from 3.87447
281/281 ————— 46s 165ms/step - accuracy: 0.9447 - loss: 0.7874
Epoch 4/30
281/281 ————— 0s 171ms/step - accuracy: 0.9574 - loss: 0.7398
Epoch 4: val_loss did not improve from 3.87447
281/281 ————— 48s 172ms/step - accuracy: 0.9574 - loss: 0.7399
Epoch 5/30
281/281 ————— 0s 163ms/step - accuracy: 0.9543 - loss: 0.7281
Epoch 5: val_loss did not improve from 3.87447
281/281 ————— 46s 164ms/step - accuracy: 0.9542 - loss: 0.7282
Epoch 6/30
281/281 ————— 0s 167ms/step - accuracy: 0.9561 - loss: 0.7039
Epoch 6: val_loss did not improve from 3.87447
281/281 ————— 47s 169ms/step - accuracy: 0.9561 - loss: 0.7039

```

```


Best accuracy: 0.9497330784797668

```

```

1 plt.plot(history.history['accuracy'], label='Training Accuracy')
2 #plt.show()

```



Conclusion:

- The model is finally up to satisfaction.
- Training and Test paths are defined (train_dir, test_dir)
- Datasets were created (train_ds, validation_ds, test_ds)
- Data for all 9 classes are visualized
- Model-1 built with vanilla dataset - accuracy was low and overfitting was noticed
- Model-2 built and run with keras augmentation - accuracy was better but still low with overfitting
- Model-3 built and run with Augmentor created dataset with steps to curtail overfitting
- Data augmentation done with Keras and Augmentor libraries
- Class distribution was done and imbalance was noticed. It is fixed with Augmentor
- Coded with nbpep8 lib and documentation where it is required.

